

Métodos Avanzados en Estadística

Estudiante

Máster en Ciencia de Datos

Universidad Autónoma de Madrid

tu-web.es



Este libro se distribuye bajo una licencia CC BY-NC-SA 4.0.

Eres libre de distribuir y adaptar el material siempre que reconozcas a los autores originales del documento, no lo utilices para fines comerciales y lo distribuyas bajo la misma licencia.

creativecommons.org/licenses/by-nc-sa/4.0/

Métodos Avanzados en Estadística

Estudiante

Máster en Ciencia de Datos

Universidad Autónoma de Madrid

tu-web.es

Índice

I. Bootstrap	5
II. PDF Estimation	6
III. Regression	7
1. Settlement	7
2. Non parametric regression. Nadaraya-Watson Estimator	8
2.1. Integrated MSE of the N-W estimator	9
2.2. Penalized MSE	10
2.3. Multiple linear regression	10
2.4. MSE fit	11
2.4.1. Variability decomposition	12
2.5. Reduced and complete models	13
2.5.1. Prediction errors	13
IV. Classification	16

Parte I.

Bootstrap

Parte II.

PDF Estimation

Parte III.

Regression

1. Settlement

The general goal in a regression problem is, given a random variable Y and a vector of random regressor variables $X = (X_1, \dots, X_p)$, study the relation between X and Y . Remark that Y can also be a random vector and in this case we are in the case of multivariate regression.

Definition 1.1. In the previous conditions, we call the regression function to:

$$m(X) = E[Y|X], \quad m(x) = E[Y|X = x].$$

It is known that $m(X)$ is the best prediction of Y from X (using the MSE). Our goal is to estimate $\hat{m}(x)$ using a finite number of i.i.d. datapoints $(x_1, y_1), \dots, (x_n, y_n)$ where $x_i \in \mathbb{R}^p$.

Usually, regression models are expressed as follows:

$$Y = m(X) + \epsilon,$$

where $E(\epsilon|X) = 0$ or, equivalently, $m(X) = E[Y|X]$. That can be proved the following way:

$$E[Y|X] = E[m(x) + \epsilon|x] = m(x) + E[\epsilon|x] = m(x)$$

because the expectation $E[\epsilon|x] = 0$. These hypothesis provide the following statements:

- $E(\epsilon) = E(E(\epsilon|X)) = E(0) = 0$
- $Var(\epsilon = 0)$ using the formula of the total variance.
- $E(\epsilon X) = E[E(\epsilon X|X)] = E[E(\epsilon|X)] = 0$.
- Lastly, the last item implies that $Cov(\epsilon, X) = E(\epsilon X) - E(\epsilon)E(X) = 0 - 0 \cdot E(X)$.

There are many different possible models:

1. Linear regression, where

$$m(X) = m(X_1, \dots, X_p) = \beta_0 + \beta_1 X_1 + \dots + \beta_p X_p$$

2. Non parametric models, that assume continuity or smoothness conditions for $m(x)$.

2. Non parametric regression. Nadaraya-Watson Estimator

Let us assume that the vector X, Y has a joint p.d.f. $f(x, y)$ and the marginal p.d.f. of X is $g(x)$. Then, using the definition of the conditional probability:

$$m(x) = E[Y|X = x] = \int y f(y|x) dy = \frac{\int y f(x, y) dy}{g(x)}.$$

A very good first idea is to use a kernel estimator in the numerator and in the denominator to replace the densities that appear in the last expression.

$$\hat{g}(x) = \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - X_i}{h}\right)$$

and

$$\hat{f}(x, y) = \frac{1}{nh^2} \sum_{i=1}^n K\left(\frac{x - X_i}{h}\right) K\left(\frac{y - Y_i}{h}\right).$$

Proposition 2.1. In the previous conditions, the Nadaraya-Watson estimator is has the expression:

$$\hat{m}(x) = \frac{\sum_{i=1}^n K\left(\frac{x - X_i}{h}\right) Y_i}{\sum_{i=1}^n K\left(\frac{x - X_i}{h}\right)}$$

Demostración.

$$\begin{aligned} \int y \hat{f}(x, y) dy &= \frac{1}{nh^2} \sum_{i=1}^n K\left(\frac{x - X_i}{h}\right) \int y K\left(\frac{y - Y_i}{h}\right) dy \\ &\stackrel{(1)}{=} \frac{1}{nh} \sum_{i=1}^n K\left(\frac{x - X_i}{h}\right) \underbrace{\int (uh + Y_i) K(u) du}_{(2)} \end{aligned}$$

where, in (1) we have applied the change of variables $u = \frac{y - Y_i}{h}$ and in (2) we have used that $\int uk(u)du = 0$ and that $\int k(u)du = 1$. To end this proof, we insert this term in the expression of $\hat{m}(x)$. \square

Proposition 2.2. The N-W estimator in x is the value $\hat{\beta}_0$ that minimizes

$$\sum_{i=1}^n w_i(x) (Y_i - \beta_0)^2,$$

which is a sum of weighted least squares.

We can change the way of expressing $\hat{m}(x)$. If we name

$$w_i(x) = \frac{K\left(\frac{x - X_i}{h}\right)}{\sum_{i=1}^n K\left(\frac{x - X_i}{h}\right)},$$

then $\hat{m}(x) = \sum w_i(x) y_i$.

2. Non parametric regression. Nadaraya-Watson Estimator

Remark 2.1. The estimator is an average of Y_i locally weighted such that, to estimate $m(x)$, we use only the Y_i such that $X_i \approx x$.

A natural question is to consider the extreme cases of the smoothing parameter h of the kernel.

- If $h \rightarrow \infty$, we have that

$$\lim_{h \rightarrow \infty} w_i(x) = \frac{K(0)}{nK(0)} \frac{1}{n}$$

and, then,

$$\hat{m}(x) = \sum \frac{1}{n} Y_i = \hat{Y}$$

- If $h \rightarrow 0$, since $K\left(\frac{x-X_i}{h}\right) \rightarrow 0$ when $j \neq i$ and $K(0)$ if $i = j$, we obtain

$$\lim_{h \rightarrow 0} \hat{m}(x_i) = Y_i$$

We can generalize the previous case as follows:

Proposition 2.3. The N-W estimator can be expressed as

$$\hat{m}(x) = \hat{\beta}_0 + \hat{\beta}_1 x + \cdots + \hat{\beta}_p x^p$$

where $\hat{\beta}_1, \dots, \hat{\beta}_p$ minimize the term:

$$\sum_{i=1}^n w_i(x) (Y_i - \beta_0 - \beta_1(x_i - x) - \cdots - \beta_p(x_i - x)^p)^2.$$

2.1. Integrated MSE of the N-W estimator

We can also approximate the bias and the variance from the NW estimator.

Proposition 2.4. We can approximate the integrated variance of the N-W estimator as

$$\int \text{Var}[\hat{m}(x)] dx \approx \frac{\sigma^2 \|K\|_2^2}{nh} \int \frac{dx}{h(x)}, \quad nh \gg 0.$$

The bias of this estimator can also be approximated as:

$$\int (E(\hat{m}(x) - m(x))^2) dx \approx \frac{h^4}{4} \sigma_K^4 \int \left(m''(x) + 2 \frac{m'(x)g'(x)}{g(x)} \right)^2 dx$$

Note that:

- Since $g(X)$ is the density of X , if $g(x) \approx 0$, we do not have much information of X around x , so $\hat{m}(x)$ has a high variance.
- In the bias, the term $\frac{2m'(x)g'(x)}{g(x)}$ is a design bias depending on \mathbf{X} . This term disappears when locally linear regression is used.

2.2. Penalized MSE

In this case, the estimators are functions that minimize $\phi(\lambda)$ for $\lambda > 0$ and

$$\phi(\lambda) := \sum_i (Y_i - m(x_i))^2 + \lambda \int_a^b m''(x)^2 dx$$

The first term measures how our model fits the data. The second term controls the smoothness of our estimator. We can think of it as the regularization term used frequently in machine learning.

We have to consider the following extreme situations:

1.

$$\lim_{\lambda \rightarrow \infty} \hat{m}(x) = \hat{\beta}_0 + \hat{\beta}_1 x$$

2.

$$\lim_{\lambda \rightarrow 0} \hat{m}(x) = \text{interpolator}$$

Remark 2.2. The solution of this problem, which is a tradeoff between fitting our data and being smooth, is a spline.

Proposition 2.5. The function that minimizes $\phi(\lambda)$ is a cubic natural spline which nodes correspond to x_1, \dots, x_n .

2.3. Multiple linear regression

The most common regression model is:

$$Y = \beta_0 + \sum_j^p \beta_j X_j + \epsilon$$

where $E(\epsilon|X_1, \dots, X_p) = 0$ and its variance is σ^2 . We can assume that, for many inferences (intervals, contrasts,...) $\epsilon|(X_1, \dots, X_p)$ follows a gaussian distribution.

Each observation from the training sample follows the model:

$$y_i = \beta_0 + \sum_j^p \beta_j x_{i,j} + \epsilon_i, \quad i = 1, \dots, n$$

where, again, $E(\epsilon_i|x_i) = 0$ and $Var(\epsilon_i|x_i) = \sigma^2$.

We can formulate the previous model as follows:

$$\underbrace{\begin{pmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{pmatrix}}_{\mathbf{Y}} = \underbrace{\begin{pmatrix} 1 & x_{1,1} & \dots & x_{1,p} \\ 1 & x_{2,1} & \dots & x_{2,p} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{n,1} & \dots & x_{n,p} \end{pmatrix}}_{\mathbf{X}} \underbrace{\begin{pmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_p \end{pmatrix}}_{\boldsymbol{\beta}} + \underbrace{\begin{pmatrix} \epsilon_1 \\ \epsilon_2 \\ \vdots \\ \epsilon_n \end{pmatrix}}_{\boldsymbol{\epsilon}}.$$

2. Non parametric regression. Nadaraya-Watson Estimator

Using the names in the under braces, our problem is formulated as:

$$Y = X\beta + \epsilon, \quad \epsilon|X \equiv N_n(0, \sigma^2 \mathbb{I}_n) \Leftrightarrow Y|X \equiv N_n(X\beta, \sigma^2 \mathbb{I}_n)$$

we call X the design matrix.

Note that, throughout this section, the data X is fixed. Hence, in the expression $Y = X\beta + \epsilon$, the term $X\beta$ is *constant*, so the distribution that Y follows depends only on the distribution of ϵ . Applying how the multi dimensional gaussian distribution is affected by sums and product by scalars, is how we obtain the distribution of $Y|X$.

2.4. MSE fit

This is the case where we want to fit our coefficients β_0, \dots, β_p using the MSE. We want to find the values β_0, \dots, β_p that minimize

$$\|Y - X\beta\|_2^2 = \sum_{i=1}^n [Y_i - (\beta_0 + \beta_1 x_{i,1} + \dots + \beta_p x_{i,p})]^2$$

It can be shown that \hat{Y} is the orthogonal projection.

Definition 2.1. In the previous conditions, the residuals vector \mathbf{e} is defined as

$$\mathbf{e} = Y - \hat{Y} = Y - X\hat{\beta}.$$

This residuals must be orthogonal to the columns of the design matrix.

$$X'(Y - \hat{Y}) = 0 \iff X'\mathbf{e} = 0$$

Proposition 2.6. The least squares estimators $\hat{\beta}$ is expressed as

$$\hat{\beta} = (X'X)^{-1}X'Y$$

To proof this, derive the error expression and then set it equal to zero to find the minimum.

Proposition 2.7. The least squares estimator is the maximum likelihood estimator of β .

$$L(\beta, \sigma^2) = \left(\frac{1}{\sigma \sqrt{2\pi}} \right)^n \exp \left\{ -\frac{1}{2\sigma^2 \|Y - X\beta\|_2^2} \right\}$$

Proposition 2.8. The vector $\hat{\beta}$ follows a $p + 1$ gaussian distribution with means vector β and covariance matrix $\sigma^2(X'X)^{-1}$.

$$\hat{\beta} \equiv N_{p+1}(\beta, \sigma^2(X'X)^{-1}).$$

If we fit the values, we obtain the adjusted values vector, defined as:

$$\hat{Y} = X\hat{\beta} = HY, \quad \text{where } H = X(X'X)^{-1}X'.$$

2. Non parametric regression. Nadaraya-Watson Estimator

\mathbf{H} is called the *hat matrix* and, geometrically, it is a projection matrix over V . Using this hat matrix, the residual vector is defined as:

$$\mathbf{e} = \mathbf{Y} - \hat{\mathbf{Y}} = (\mathbf{I} - \mathbf{H})\mathbf{Y}$$

We also want to estimate the **variance** σ^2 . To do this, we can use the variance of our residuals:

$$\mathbf{S}_R^2 = \frac{1}{n-p-1} \sum_{i=1}^n \mathbf{e}_i^2$$

Proposition 2.9.

$$(n-p-1)\mathbf{S}_R^2/\sigma^2 \equiv \chi_{n-p-1}^2$$

This result is a consequence of the fact that the distribution of the residuals follows a gaussian distribution, and we are multiplying by a constant that normalizes each term, so we are adding Gaussian distributions, which is the definition of a χ^2

Note. Thanks to the previous proposition, we can build confidence intervals and contrasts for σ^2

Proposition 2.10. \mathbf{S}_R^2 and $\hat{\beta}$ are independent.

2.4.1. Variability decomposition

Definition 2.2. We can define the following quantities:

- The **Total squared sum**

$$SCT = \sum_{i=1}^n (\mathbf{Y}_i - \bar{\mathbf{Y}})^2$$

measures the total variability of the predicted variable.

- The **Explained squared sum**

$$SCE = \sum_{i=1}^n (\hat{\mathbf{Y}}_i - \bar{\mathbf{Y}})^2$$

measures the variability of the predicted variable *explained by the model*.

- The **Residual squared sum**

$$SCR = \sum_{i=1}^n e_i^2$$

measures the variability of the predicted variable that is *not explained by the model*.

2. Non parametric regression. Nadaraya-Watson Estimator

Using that the residuals are orthogonal to the regression variables, we obtain

Corollary 2.1. The total variability of the predicted variable can be decomposed into the explained variability by the model and the residual variability (not explained by the model). That is:

$$SCT = SCE + SCR$$

Definition 2.3. The determination coefficient measures the capability of our model to explain Y .

$$R^2 = \frac{SCE}{SCT}$$

Let us set in the case that we would like to contrast the hypothesis

$$H_0 : \beta_1 = \dots = \beta_p = 0.$$

In this case, we would use the statistic:

$$F = \frac{SCE/p}{SCR/(n-p-1)}.$$

Under H_0 , the statistic F follows a distribution $F_{p,n-p-1}$.

2.5. Reduced and complete models

Adding complexity to a model usually turns into a better fit to available data. However, this could result in worse predictions for new data. The simpler the model, the less overfitting.

Our goal in this section will be to compare two models:

1. The complete model
2. A simplified model M_0 , where we constrain our model to $A\beta = 0$, where $A \in \mathcal{M}_{k \times (p+1)}$ and $range(A) = k < p + 1$.

We will compare both models contrasting $H_0 : A\beta = 0$.

2.5.1. Prediction errors

Suppose that we have a real model, predicting some data. Let $Y = (Y_1, \dots, Y_n)$ be the labels of the training data, $\tilde{Y} = (\tilde{Y}_1, \dots, \tilde{Y}_n)$ the labels of the test data. We hypothesize that the vectors Y, \tilde{Y} are independent. We assume that $E(Y) = E(\tilde{Y}) = \mu = (\mu_1, \dots, \mu_n)$ and Σ be the covariance matrix.

We fit a regression model $Y = X\beta + \epsilon \iff Y = \mu + \epsilon$ where $\mu \in V = \{X\beta : \beta \in \mathbb{R}^p\}$. We already know that:

- The MSE model is $\hat{\beta} = (X'X)^{-1}X'Y$
- ... D(44/57)

Definition 2.4. We call the training error:

$$E_{\text{train}} E \left[\sum_{i=1}^n (Y_i - x_i' \hat{\beta})^2 \right]$$

and we can define the test error the same way using test labels.

Proposition 2.11. The test error can be expressed as:

$$E_{\text{test}} = E_{\text{train}} + 2E(\|C\|^2)$$

Demostración. Recall that:

$$E_{\text{test}} = E(\|A\|^2) + E(\|B\|^2) + E(\|C\|^2) + 2E(A'B) + 2E(A'C) + 2E(B'C)$$

and, since $E(A'B) = E(A'C) = E(B'C) = 0$ and using that $E(\|B\|^2)$ and $E(\|C\|^2)$ are identical and iid, we obtain the final result. \square

Having the previous proposition, we would like to calculate the last term of the error:

$$E(\|C\|^2) = E[(\hat{\beta} - \beta)' X' X (\hat{\beta} - \beta^*)]$$

We can use the following *trick*:

$$E(\|C\|^2) = E[(\hat{\beta} - \beta)' X' X (\hat{\beta} - \beta^*)] = E(\|C\|^2) = E[\text{tr}((\hat{\beta} - \beta)' X' X (\hat{\beta} - \beta^*))] = E(\|C\|^2) = E[\text{tr}(X X' (\hat{\beta} - \beta^*))]$$

and, recalling that $X(X'X)^{-1}X' = H$, we obtain that

$$\text{tr}(E[(X'X)(X'X)^{-1}X'\Sigma X(X'X)^{-1}]) = \text{tr}(\Sigma H).$$

We have just proved that

$$E_{\text{test}} = E_{\text{train}} + 2\text{tr}(\Sigma H).$$

Furthermore, recalling that

$$\text{Cov}(Y, \hat{Y}) = \text{Cov}(Y, HY) = \Sigma H$$

we obtain that

$$E_{\text{test}} = E_{\text{train}} + 2 \sum_{i=1}^n \text{Cov}(Y_i, \hat{Y}_i).$$

An interpretation for this is that, if $Y_i \approx \hat{Y}_i$, then we are being pretty optimistic.

Case: Independence and homocedasticity

Consider that

$$\begin{pmatrix} \sigma^2 & 0 \\ 0 & \sigma^2 \end{pmatrix} = \sigma^2 I$$

then,

$$\text{tr}(\Sigma H) = \sigma^2 \text{tr}(H) = \sigma^2 p$$

2. Non parametric regression. Nadaraya-Watson Estimator

$$\text{tr}(X(X'X^{-1}X')) = p$$

Hence,

$$E_{\text{test}} = E_{\text{train}} + 2p\sigma^2$$

Case: Independence and heterocedastic XDXXD

In this case,

$$\begin{pmatrix} \sigma_1^2 & 0 \\ 0 & \sigma_n^2 \end{pmatrix} = \sigma^2 I$$

Hence,

$$E_{\text{test}} = E_{\text{train}} + 2\sigma_i^n \sigma_i^2 h_{ii}$$

Parte IV.

Classification